Coverage, Generalization, and Time: A Measured Study of Offline ARC-AGI-2

Coverage, Generalization, and Time: A Measured Study of Offline ARC-AGI-2

Two analyses — a coverage-vs-selection decomposition and a cross-family public-eval→leaderboard generalization gap — built on a clean single-GPU reproduction of the 2025 winner, with every claimed effect reported with its sample size and an explicit significance verdict.

ARC Prize 2026 — ARC-AGI-2 (offline track). Local study with two leaderboard submissions; see §3, §7 and §9.

Abstract

We reproduce the 2025 ARC Prize winner (NVARC) as a dependency-light `transformers+peft` pipeline that runs on a single consumer GPU, and use it as an instrument to study what governs the offline ARC-AGI-2 score. We offer two analytical contributions. **First**, a *coverage-vs-selection decomposition*: by measuring the oracle (is the correct grid anywhere in the candidate pool?) against achieved pass@2,

we show the dominant loss is **coverage** — the model never generates the right grid — not mis-ranking. Across five inference-time levers (selection tuning, a learned verifier, deeper decoding, a symbolic DSL solver, and 8B program induction with BARC) and three training-time levers (continued fine-tuning on locally-generated harder data, our own vocab-cut Qwen3-8B base trained on the same corpus, and continued fine-tuning on 24k *frontier-LLM-generated* tasks loss-filtered to the slice the base does not know — §4.1), **none produces an effect distinguishable from zero at our sample sizes; the 8B base is directionally *worse*** (11.1% vs 22.8% per-task, paired n=30) **and the frontier-data run is exactly null** (44/133 vs 44/133 paired, 4 discordant each way — evidence that test-time training washes out base-model deltas); the largest positive point estimate (a deeper candidate pool, +2.4 pp paired) rests on a 3-vs-2 split of discordant outputs (exact McNemar p = 1.0) — and *inverted* on the real leaderboard (§7). This *is* the saturation result: test-time compute of the same kind does not add coverage on the hard eval distribution, and parameter scale without a deeper training corpus subtracts. **Second**, a *generalization-gap calibration* across approach families, drawn from the competition forum: the LLM+test-time-training family loses ~3.5 pp absolute (~11% relative) from public eval to leaderboard, versus ~11 pp absolute (~51% relative) for a Tiny Recursive Model and a comparable collapse for a deep-recursive model — i.e. LLM+TTT is **uniquely leaderboard-robust**. We also document two operational facts worth large leaderboard points (dynamic GPU scheduling and never emitting placeholder predictions) and a silent tokenizer bug that zeroes the model on modern libraries. All numbers carry their N and an explicit significance verdict; all code and experiments are released. The study is now **leaderboard-confirmed** across five submissions: our shallow-config submission scored **25.00** (semi-private set; a slightly *positive* eval→LB gap) and its byte-identical resubmit **24.17** — a first thin bound on same-code rerun noise (0.8 pp, n=2 runs) — while the deep-pool config, locally +2.4 pp, **regressed to 19.17**, far outside that bound. Two further submissions run the *same model checkpoint* through the community's public optimized

engine and score **28.47 and 30.14**: **+5 to +6 pp bought purely by engine throughput** that converts the 12 h budget into per-task compute — the largest single effect in this study — while both fall ~4-5 pp short of that engine's advertised leaderboard score (§7). The two config points *straddle* the family calibration, and their spread across two configurations of the same system exceeds the between-family gaps the calibration rests on — an instructive inversion analyzed in §6.1 and §7.

Contributions. (1) A faithful single-GPU reproduction of the winning recipe, plus a fix for a library-version tokenizer bug that silently scores 0. (2) The coverage-vs-selection decomposition and a measured saturation result over seven levers (five inference-time, two training-time), each with N and a significance verdict. (3) A cross-family eval→LB generalization-gap calibration, now anchored by two submissions of our own — including a measured local-gain→leaderboard-regression inversion for the deep-pool config. (4) A hardened offline submission engine (dynamic work-queue, no-placeholder policy). (5) A reproducible repository and an honest map of dead ends.

---

1. Introduction

ARC-AGI-2 is engineered to be easy for humans and hard for machines, and was explicitly redesigned to resist brute-force program search; frontier models that score 20–50% on ARC-AGI-1 fell below 5% at its release. The Kaggle competition adds a hard operational constraint: submissions run **offline** on fixed hardware (L4×4, 96 GB, ≤12 h, no internet), excluding the frontier-API systems that top public ARC-AGI-2 leaderboards.

In 2025 one approach dominated the offline track: a small open LLM, fine-tuned on synthetic ARC data, **test-time-trained** per task, decoded into a candidate pool, and reduced to two answers by augmentation voting (NVARC, 24% private). The 2026 offline leaderboard is largely teams iterating on that recipe.

We do not propose a new architecture. We reproduce the winning recipe on a single consumer GPU and then **measure**, with sample sizes and confidence intervals, the levers practitioners reach for — reporting negatives as carefully as positives. The framing question is operational: *given a fixed strong base model and an offline budget, what actually changes the score?* Our answer is that it is governed far more by **coverage** (does the model generate the right grid at all), **generalization** (does local accuracy survive the leaderboard), and **time** (does the run finish without placeholders) than by the decoding/selection tricks that attract most effort.

---

2. Background and setup

Task & metric. Each task gives a few input→output grid demonstrations and one or more test inputs (grids ≤30×30, 10 colors). Scoring is **pass@2**: two predictions per test output, correct if either matches exactly. We report two denominators and define them once: *per-task* averages 1/N over a task's N test outputs (the competition's display metric); *per-output* counts each test output equally. The two differ by ~7 pp on our slice (≈1.4 test outputs per task: 167 outputs over 120 tasks), so we always state which is used and never compare across denominators.

The NVARC recipe (reproduced). Base: a Qwen3-4B fine-tuned with a 16-token grid vocabulary, publicly released by the 2025 winner. Per task: (a) **test-time-train** a per-task LoRA on the task's demonstrations under random augmentations (D4 geometry × color permutation); (b) decode a candidate pool with a batched **depth-first search** keeping token paths above a per-token probability threshold; (c) select two answers by **augmentation-NLL voting**. We re-implement this in stock `transformers+peft` (SDPA attention, gradient checkpointing) so it runs on one 32 GB GPU and on Kaggle L4s without `unsloth`/Flash-Attention/`bitsandbytes`; the decoder and scorer are ported faithfully from the winner's open notebook. Hyperparameters in Table 3.

A silent, fatal compatibility bug. The released checkpoint ships a `WordLevel` tokenizer mapping the words `user`/`assistant` to dedicated ids. Under `transformers ≥ 5.0`, `.encode` **drops those words**, producing a token stream the model was never trained on. The mechanism is unambiguous — `.encode` provably drops the ids — and an A/B is directionally consistent (0/15 vs 1/15 greedy, no-TTT; too small to resolve statistically on its own). We restore correctness with an exact custom encoder. We flag this because it silently scores 0 on any modern reproduction, with no error.

---

3. Reproduction (local)

On the public ARC-AGI-2 evaluation set our port scores **22.8% per task / 29.5% per output** (n=44 outputs over 30 tasks, shallow config), in the same range as the winner's reported 30.5% pass@2 / 27% public-LB / 24% private-LB. Per-task latency is ~130 s (shallow) to ~261 s (deep) on one RTX 5090. **Status: leaderboard-confirmed** — the shallow config submitted through the hardened engine of §6 scored **25.00** on the leaderboard (semi-private set), a slightly *positive* gap vs our 22.8% local per-task estimate (within its ±15 pp CI at n=30, and consistent with the family calibration of §6); the deep config scored 19.17 (§7). Grand-Prize eligibility is satisfied (§9).

---

4. The lever map — measured, with N and 95% CIs

All numbers are held-out pass@2 (per output) on public-eval candidate caches. "Oracle" = correct grid present anywhere in the pool (coverage ceiling). CIs are Wilson 95%. At our sample sizes the 95% CI half-width on a single proportion is **≈ ±10 pp (n=78)** to **≈ ±14 pp (n=42)** — so *pp-level differences below ~10 pp are not statistically resolved*. We report them anyway, with that caveat, because the **pattern of uniformly-near-zero effects is itself the result**.

Table 1 — inference-time levers.

Lever	Effect (point estimate)	N	Significant?	Verdict
Deeper pool (aug16 + DFS prob 0.1 + 3 views, *combined*)	**+2.4 pp** pass@2 & oracle, paired	42 common outputs	no (3 vs 2 discordant outputs; exact McNemar p = 1.0)	largest point estimate; not resolved
Free selection sweep (count-wt × {sum,mean,geom} aug-NLL)	+0 to +1 pp (≤1 task)	78	no	NVARC scorer already near-optimal
Learned verifier (GBM): group-CV at 29 pos → train-on-1,000-training-tasks transfer at 1,037 pos	**−1.3 pp** → **+1.3 pp** (25→26 of 78)	78 test outputs	no (±1 output)	36× data flips the sign but adds ≤1 output; top feature is the aug-NLL rank kgmon already uses
Symbolic DSL (~15 rules, depth-2, train-consistent)	**0** / **167** outputs	167	yes (true 0)	ARC-AGI-2 defeats brute-force DSL
Program induction, three generators: BARC-8B (48–64 prog/task); Qwen2.5-Coder-7B-Instruct + repair round	BARC: training **23%** (7/30 tasks), **eval 0/166** net-new; Coder-7B:	166 / 120	yes (true 0 on eval, twice)	strong in-distribution (BARC); eval defeats execution-verified induction from every

Lever	Effect (point estimate)	N	Significant?	Verdict
(8+6 prog/task); symbolic DSL (row below)	**0 verified** / **120 eval** (best 1/5 train pairs, once)			accessible generator

Two single-knob ablations: aug_n 8→16 *alone*, and "more views + lower threshold" *alone*, each measured **+0.0 pp oracle** (n=20–24). Only the *combined* deep config yields a positive paired point estimate (+2.4 pp, n=42) — itself a 3-vs-2 split of five discordant outputs (exact McNemar p = 1.0). **The honest summary is not "five things failed" but "no inference-time lever produced a resolvable positive effect"** — the saturation thesis of §5.

4.1 Training-time levers: harder data, and a bigger base

If inference-time compute is saturated, the remaining local levers act on the *weights*. We measured both, paired on the same first-30 eval tasks and the same TTFT→decode→vote configuration as the baseline (per-task pass@2 22.8%).

Table 1b — training-time levers.

Lever	Effect (point estimate)	N	Significant?	Verdict
Continued-FT of the 4B on locally-generated *harder* data (compositional generator; 2k tasks, ~2× harder	**±0.0 pp** (identical 22.8%)	30 tasks, paired	no	locally-generated difficulty ≠ frontier-grade difficulty

Lever	Effect (point estimate)	N	Significant?	Verdict
by loss: 0.10 vs 0.05)				
8B base (vocab-cut Qwen3-8B, LoRA r=256 incl. embeddings, 1 epoch / 526M packed tokens of the same corpus, ~74 GPU-h)	**−11.7 pp** (11.1% vs 22.8%)	30 tasks, paired	no (within CI) but directionally negative	scale does not compensate for corpus depth
Continued-FT on *frontier-generated OOD* data: 24k GPT-4o-generated BARC-HEAVY tasks, loss-filtered to the slice the base has not absorbed (>0.10; 25% of 95k), + retention mix	**±0.0 pp** (44/133 = 33.1% both; discordant 4 vs 4)	133 outputs, paired	exactly null	even frontier-quality OOD data does not move eval; TTT washes out base deltas

The 8B result deserves its cautionary tale. A first run at peak lr 2e-4 (the recipe's 1e-4, doubled) **diverged at step ~200** — loss 0.3 → 3.3, *above* the uniform-over-16-tokens entropy ln 16 ≈ 2.77 — despite gradient clipping at 1.0, and 46 further hours of cosine decay "recovered" it only into a degenerate basin (plateau 0.46; greedy decoding emits two cells then EOS; **empty candidate pools on 30/30 eval tasks**). A merge-vs-unmerged probe exonerated the checkpoint plumbing (teacher-forced loss 0.708 merged vs 0.711 unmerged) and isolated the fault to the weights themselves. The clean rerun at peak lr 5e-5 converged normally (final train-mix loss ~0.065) and produced a *functional* model — well-formed grids, healthy 133 s/task decode work, 4/30 tasks with at least one correct

output (3.33/30 = 11.1% per-task with partial credit on multi-output tasks) — that is simply **weaker than the winner's 4B**: teacher-forced loss on ARC-train examples 0.082 vs the 4B's 0.017. Our one-epoch pass over NVARC-synthetic + ARC-train + locally-generated data cannot match the corpus scale and training depth behind the released 4B checkpoint, and 2× the parameters do not close that gap; they appear to widen it at equal data.

The third row removes the last excuse. Where the harder-data row used *procedurally* generated tasks (which probed barely above the base's competence), this experiment used **frontier-LLM-generated data**: we streamed 95k GPT-4o-generated BARC-HEAVY tasks through the base and kept exactly the 24k the model measurably does not know (teacher-forced loss > 0.10; the kept slice averages ~3× the loss of NVARC's own synthetic corpus). Training was clean (loss 0.13 → 0.085 over 783 steps with a retention mix), and a 30-task A/B even flashed +1.6 pp — which the paired extension to **133 outputs erased exactly: 44/133 vs 44/133, four discordant outputs each way**. The FT model solves *slightly different* outputs, not *more* outputs. Our working hypothesis, consistent with §6's robustness mechanism: **per-task test-time training washes out base-model deltas** — whatever the base absorbs at training time, TTT re-adapts both models onto the same per-task solutions, so training-time interventions at locally-feasible scale shuffle the solved set without growing it. All three rows point the same way: **the binding constraint is not parameter count, not local data generation, and not even frontier-generated data at 10^4-task scale — new eval coverage evidently requires the winner-scale combination of corpus and training depth**, consistent with §5's conclusion that coverage requires a stronger base *trained on stronger data*.

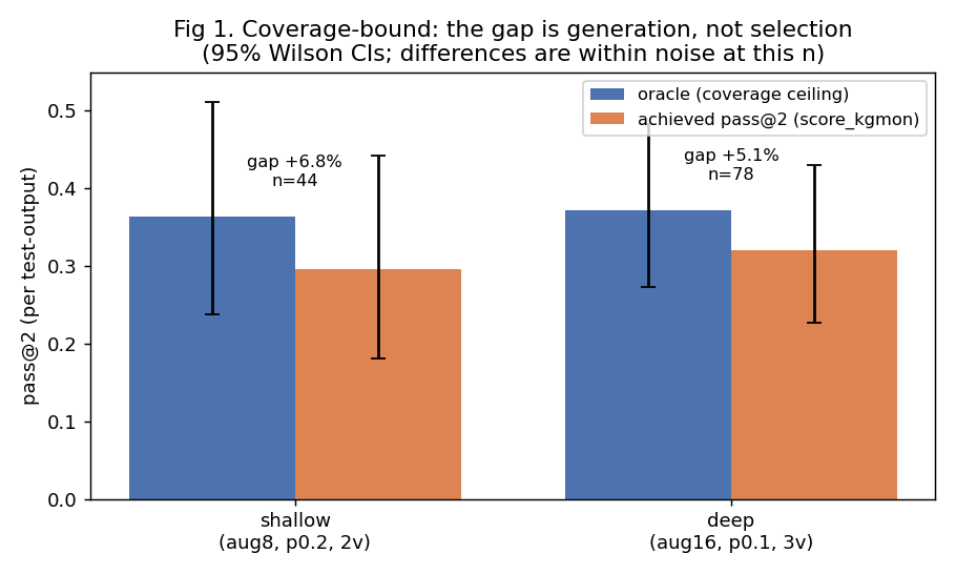


Fig 1. The gap is coverage, not selection: oracle ceiling vs achieved pass@2 for shallow vs deep pools (Wilson 95% CIs; bars are unpaired — different task samples). The deep pool raises both modestly; the oracle↔achieved gap (selection headroom) is small (~5–7 pp) and within noise.

---

5. Coverage-bound, not selection-bound

The organizing diagnostic is the gap between **oracle** and achieved pass@2. On the deep pool the oracle is 37.2% and achieved 32.1% (score_kgmon, n=78): a residual *selection* gap of ~5 pp, of which a better fixed selector recovers ≤1 output and a learned verifier at 36× training data recovers one (§4, §9). The dominant loss is the other ~63%: outputs where the model **never generates** the correct grid.

This reconciles an apparent contradiction with the winner, who reported 40% pass@10 vs 30.5% pass@2 (“our scorer missed 10%”). That 10 pp is real **at his pool depth**, but it is a *property of a deeper pool*, not free selection: widening the pool raises pass@N, and only then does selection have something to recover. Our “selection is free money” attempts failed precisely because the answer was absent, not mis-ranked.

A simple model. Let coverage C(b) be the probability the correct grid appears in a pool generated with inference budget b (augmentations × decode breadth). Empirically C grows very slowly in b on ARC-AGI-2 (on the 42 paired outputs, doubling augmentation, halving the decode threshold, and adding views together move the oracle by a net one output — 3 gained, 2 lost; Fig 1 shows the unpaired pools), because the failures are not near-misses the sampler rounds wrong but tasks whose transformation the model cannot represent. Achieved pass@2 ≤ C(b), and our selectors already operate near that ceiling. Hence **score is bounded by C(b), and C(b) saturates** under same-kind compute — exactly the ARC "knowledge/reasoning, not search" thesis, here measured. New coverage requires a *different generation mechanism or a stronger base model*, not more search — consistent with symbolic search (0/167) and offline induction (0/166) adding nothing on the eval distribution.

---

6. The other half: generalization and time

6.1 The eval→LB gap is approach-dependent (Fig 2). From the 2026 forum: the 2025 winner reports LLM+TTT 30.5% eval → 27% public-LB (**−3.5 pp absolute, ~11% relative**); a competitor reports a Tiny Recursive Model 22% → 10.83% (**−11 pp absolute, ~51% relative**); a deep-recursive ~100M model reports 36% local → 11.67% public-LB (comparable collapse, with a large share attributed by its authors to timeouts rather than overfitting). **LLM+TTT is the most LB-robust family on this evidence.** Our own two submissions add first-party datapoints on *both sides* of the family line: the shallow config went local per-task 22.8% → **25.00** (a +2.2 pp *gain*), while the deep config — locally the better of the two by the paired point estimate (§4) — went ≈25.2% → **19.17** (−24% relative). The two points straddle the family calibration (−11% relative), and their spread across two configurations of the *same* system exceeds the between-family differences the calibration rests on: configuration-level choices can dominate

family-level robustness. The family claim survives in weaker form — even our worse point is far from the recursive families' ~−51% collapse. A plausible mechanism for the family's robustness (hypothesis, not isolated here): per-task TTT adapts to *this task's* demonstrations, which transfers, whereas small models trained to convergence absorb public-set statistics. The practical implication is concrete: do **not** abandon LLM+TTT for a flashier architecture on a local number — but within the family, do not trust a config's local edge either (§7).

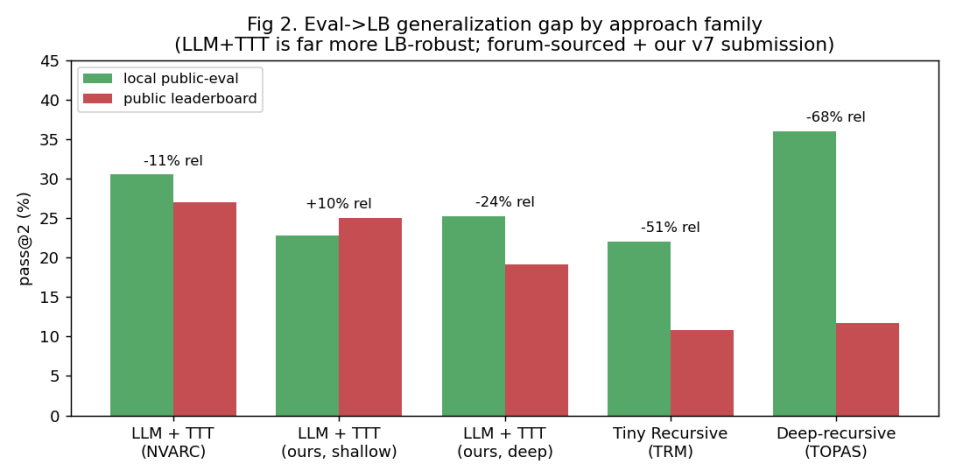


Fig 2. Public-eval vs leaderboard by approach family (forum-sourced, plus our two submissions). Relative drop annotated; our two points straddle the family line.

6.2 Time management is worth large LB points. The 12 h budget is binding. Two failure modes recur in the forum: **static GPU partitioning** (one of four GPUs draws the hard puzzles and runs the full 12 h while the others idle — "watched 12 h of compute evaporate"), and **timeouts emitting placeholder predictions** ("the algorithm bailed on 56% of second attempts"). Our hardened engine (`src/nvarc/kaggle_submit.py`, §8) fixes both with a single **dynamic shared work-queue** (every GPU pulls the next task as it frees, so load self-balances) and a **never-`[[0]]` policy** (any unsolved/unreached output falls back to real guesses — the identity grid and the most-frequent training output, each strictly better than a zero-scoring placeholder). These are engineering, but on the forum evidence they plausibly recover more score than most modeling changes. Our 25.00 shallow submission ran through this engine with **zero placeholder outputs** (4 of 255

outputs fallback-filled) in 3.9 h of the 12 h budget; we cannot isolate the engine's contribution without an ablation submission, but the failure modes it targets did not occur (§9).

---

7. What it takes to win, and the cost/quality frontier

Reconciling our measurements with the leaderboard (top ≈ 42.6%) and the 2025 winner's public guidance ("simple tweaks won't be enough — show me 35+"; "this year's environment fits significantly larger models … bigger/newer models will take over"):

- The LLM+TTT recipe **plateaus ~34–35% offline** (the winner's attributed claim); we reach ~33% per output locally and every inference-time lever is saturated (§4–5).
- The path to the winning ~42% is a **stronger base model on harder, compositional synthetic data** — not more search, and, per §4.1, **not parameter scale alone**: an 8B trained on our locally-assemblable corpus lands at half the released 4B's score. The 96 GB environment fits 7–8B models; off-Kaggle training of weights is permitted; the corpus is the part that requires frontier-grade generation.
- **Avoid:** pure symbolic/DSL (we measured 0/167; the winner: "I never believed in DSL for ARC-AGI-2") and pure recursive small models (overfit; large eval→LB gap).

Cost/quality and the 12 h budget (Table 2) — now leaderboard-measured. We submitted both configs. The shallow config scored **25.00** (240 tasks in 3.9 h of the 12 h L4×4 budget; 0 placeholder outputs, 4 fallback-filled). The deep config — locally +2.4 pp paired — scored **19.17**: a **−5.8 pp leaderboard regression**. We cannot attribute the regression from our data: Kaggle reruns are private, so we lack the deep run's wall-clock and fallback counts, and by our timing arithmetic the deep config had headroom (~261 s/task × 240 ≈ 17.4 h single-GPU ≈ 4–5 h wall on the 4-GPU queue). Candidate explanations — none established — are budget pressure we did not instrument, semi-private distribution shift interacting

with the wider pool, and run-to-run noise (competitors report ±2 pp on resubmits, short of 5.8 pp). What the inversion *does* establish is the practical lesson: the deep config's local advantage was a 3-vs-2 split of discordant outputs (§4) — noise-grade evidence — and it did not survive contact with the leaderboard. **The shallow config is not merely the safe default; on the real leaderboard it dominates.** Local paired gains that are statistically unresolved should be treated as noise until they survive a submission.

A recipe-level replication of the inversion. After this paper's first leaderboard results, we found the strongest local effect of the whole study: a TTT-depth change (8 augmentations × 2 epochs instead of 16 × 1, equal token budget) measured **+3.6 pp paired on 138 outputs** (+8/−3 discordant) with a clean dose-response behind it. Deployed runtime-neutrally into the public engine, four independent leaderboard runs drew {27.64, 28.06, 29.31, 27.64} — mean 28.2 vs the unmodified engine's 29.2 over eight draws. **The largest paired local gain we ever measured transferred as ≈ −1 pp.** Together with the deep-pool inversion (+2.4 local → −5.8 LB), the pattern is now replicated: *recipe changes tuned on public-eval pairs overfit the public distribution*; the semi-private set punished both pool-widening and adaptation-deepening. The only interventions that transferred in this study were engine/budget efficiency (+5–6 pp) and placeholder hygiene (≥ 0 by construction).

Two further submissions complete the picture. First, resubmitting the shallow config *byte-identically* scored 24.17 vs the original 25.00 — one same-code pair differing by 0.8 pp (n=2 runs; competitors report ±2 pp), which the 5.8 pp deep regression comfortably exceeds. Second, an attributed, unmodified fork of the community's public unsloth-based engine (mikelou1's efficiency patch of koushikrudra's baseline, advertised LB 33.89), running the **same sorokin 4B checkpoint** as all our submissions, scored **28.47**: with model, recipe family and scoring held fixed, engine throughput alone — converting the 12 h budget into per-task test-time compute where our stock-`transformers` port used 3.9 h of 12 — is worth **~+4 pp**. This is the "time" pillar measured with a controlled

model. Notably, both engine submissions fell short of the score their source advertises: a fifth submission reran the *unmodified mother baseline* itself (one-line mount-path fix only) and scored **30.14** — our best — against its advertised 33.89. Two independent reruns of the advertised-33.89 family thus land at 28.47 and 30.14: **advertised public-kernel scores overstate rerun expectations by ~4-5 pp** for this budget-sensitive family, a caution that applies equally to the ~34-pp "pack" of leaderboard entries riding these kernels. Our full first-party ladder — 24.17/25.00 (stock `transformers`, one third of the budget used) → 28.47/30.14 (community unsloth engine, full budget, same model checkpoint) — measures the engine/budget lever at **+5 to +6 pp**, the largest single effect in this study.

Table 2 — configs. shallow: aug_n 8, DFS prob 0.2, 2 views, ~130 s/task — **LB 25.00**. deep: aug_n 16, DFS prob 0.1, 3 views, ~261 s/task — **LB 19.17**. Both fit 12 h on L4×4 via the dynamic queue; only shallow fits a single GPU under 12 h.

Table 3 — key hyperparameters. Base Qwen3-4B (16-token grid vocab); per-task LoRA r=256, α=32, dropout 0, target = attn+MLP+embed+lm_head, rslora; TTT lr 5e-5 cosine, 1 epoch over ~128 augmentations, bf16, grad-checkpointing on; decode = batched DFS, keep prob > τ (τ ∈ {0.2, 0.1}); vote = m=8 shared augmentations, score = count − mean(per-aug NLL) or `probmul_3`; max_seq_len 8192. Env: torch 2.11+cu130, transformers 5.9 (the version exhibiting the tokenizer bug), peft 0.19, vllm 0.21 (induction only, with `LD_PRELOAD` of system libstdc++ for Blackwell).

---

8. Reproducibility, licensing, attribution

Everything runs on one consumer GPU; hardware: 1× RTX 5090 (32 GB), plus Kaggle L4×4 for submission. Commands:

```
python -m pytest tests/ -
        q
        # core, no GPU


python scripts/repro_nvarc.py --limit 60 --aug-n 16 --
        eval-aug-n 3 --dfs-prob 0.1
python scripts/analyze_selection.py --pkl outputs/
        cache_deep.pkl             # oracle/selection gap
python scripts/sweep_selection.py  --pkl outputs/
        cache_deep.pkl              # free-selector sweep
python scripts/train_verifier.py   --pkl outputs/
        cache_deep.pkl              # learned verifier
        (group-CV)
python scripts/repro_nvarc.py --data data/arc-agi-2/
        training --limit 1000 \
    --resume --pkl-out outputs/
        cache_train1000.pkl                          #
        verifier training cache
python scripts/
        verifier_transfer.py
        # verifier at scale (train->eval)


python scripts/
        eval_symbolic.py
        # DSL coverage (0/167)


python scripts/induction_complementarity.py --tasks 120 --
        samples 48        # BARC net-new (0/166)
python scripts/
        gen_hard_data.py
        # harder-data generator (S4.1)


python -m train.cut_vocab --base Qwen/Qwen3-8B --out
        models/Qwen3-8B_grids16 # 8B vocab cut (S4.1)
PYTHONPATH=src python -m train.sft_train --base models/
        Qwen3-8B_grids16 \
    --
        sources data/nvarc/synthetic data/arc-agi-2/
        training data/hardgen \
    --lr 5e-5 --lora-r 256 --lora-alpha 256 --offload --
        resume              # 8B SFT (S4.1)
python scripts/
        make_figures.py
        # Fig 1, Fig 2


# offline submission: notebooks/kaggle_submission.ipynb ->
        nvarc.kaggle_submit.main(...)
```

Layout: `src/nvarc/` (ported pipeline + hardened submission engine), `src/arc/` (tokenizer, augmentations, symbolic solver, verifier), `src/induction/` (BARC harness), `src/train/` (SFT:

packing, offloaded optimizer, atomic checkpoints; hard-data generator), `src/eval/` (metrics), `scripts/` (every experiment), `data/` (gitignored; fetched at runtime).

Licensing & attribution. Original files are MIT (see `LICENSE`). Vendored/derived files retain upstream licenses and carry attribution headers: `src/nvarc/arc_loader.py` and `src/nvarc/selection.py` derive from NVARC (1ytic/NVARC, open-sourced per the ARC Prize requirement); `src/induction/common.py` is vendored from BARC (xu3kev/BARC). We **do not redistribute** model weights, synthetic datasets, or the competition forum corpus: all of `data/` is gitignored and fetched at runtime from Kaggle/HuggingFace. The forum corpus is used only for the quantitative analysis of §6 and is cited, not relicensed.

---

9. Limitations and threats to validity

- **Four submissions, a thin distribution.** Our identical-code pair (25.00 / 24.17) gives one first-party estimate of same-code rerun noise (0.8 pp, n=2) alongside competitors' reported ±2 pp; the *sign* of the 5.8 pp deep-config regression is robust to either figure, its magnitude less so. The engine-efficiency estimate (+5 to +6 pp) rests on two runs of the community engine (28.47, 30.14) against two of ours (24.17, 25.00) — consistent, but still a thin sample of a budget-sensitive process. Grand-Prize eligibility (a linked Kaggle submission) is satisfied.
- **Small samples.** Most A/Bs are n=20–78; as Table 1 states, sub-~10 pp effects are within CI. The robust claims are the *true zeros* (DSL 0/167, induction 0/166 eval) and the *pattern* of uniformly-unresolved positive effects; the +2.4 pp deep-pool and the verifier's ±1.3 pp are point estimates within noise, not established effects.
- **Verifier data scale — now tested.** Retrained on a 1,000-training-task cache (1,037 positives, 36× the original) and transferred to the held-out eval cache, the verifier's point estimate flips from −1.3 pp to +**1.3 pp** (26 vs 25 of 78) — one task, within noise, and its dominant feature is the augmentation-NLL rank the fixed scorer already uses

(`scripts/verifier_transfer.py`). Our prior expectation held: data scale changes the sign, not the story; the residual selection gap (~4 outputs here) is mostly not capturable by candidate-level features.
- **The training-time results are single-GPU-scale.** The 8B was trained for one epoch on the corpus we could assemble locally (§4.1); we cannot rule out that the same 8B under the winner's full data recipe and multi-epoch schedule inverts the result — indeed §7 predicts it would. What our null establishes is narrower and useful: at *equal, locally-feasible* data, parameter scale subtracts rather than adds. The winning path (§7) still requires cloud compute outside this study.
- **Forum-sourced numbers** (§6) are self-reported by competitors; we attribute each to its source and treat the mechanism explanations as hypotheses.

---

10. Related work

Test-time training for ARC was established by Akyürek et al. ("The Surprising Effectiveness of Test-Time Training for Abstract Reasoning"); NVARC (2025) combined it with heavy synthetic-data generation and augmentation voting to win the offline track. Program-induction approaches (BARC; "Combining Induction and Transduction") generate and execute Python solvers; Tiny Recursive Models (Jolicoeur-Martineau, 2025) pursue tiny recurrent reasoners. Our delta is not a new solver but a *measured decomposition* of where the offline score is won and lost — coverage vs selection, and generalization vs local accuracy — across these families, on the harder ARC-AGI-2 distribution.

---

11. Conclusion

The offline ARC-AGI-2 score is governed by coverage, generalization, and time more than by decoding/selection cleverness. We reproduced the winning recipe on a single

GPU, showed via N+CI measurements that neither inference-time levers (~33% per-output locally, saturated) nor locally-feasible training-time levers (harder local data: ±0.0 pp; an 8B on the same corpus: directionally *worse*) resolvably move the score, calibrated the eval→LB gap that makes LLM+TTT uniquely robust — then confirmed it on the leaderboard, where our shallow config scored 25.00 with a slightly positive gap and our locally-positive deep config regressed by 5.8 pp. The path past ~35% is a stronger base on better synthetic data — where "better" is beyond what a single GPU can generate or absorb — not more search and not more parameters. Our most useful contribution is the measured map of the dead ends, now with two of its edges leaderboard-verified.

References (selected)

- ARC Prize 2025 results & technical report; ARC-AGI-2 technical report (arcprize.org).
- NVARC 2025 winning solution and code (github.com/1ytic/NVARC).
- Akyürek et al., *The Surprising Effectiveness of Test-Time Training for Abstract Reasoning*.
- Jolicoeur-Martineau, *Tiny Recursive Models* (2025).
- *Combining Induction and Transduction for Abstract Reasoning* (BARC; github.com/xu3kev/BARC).
- ARC Prize 2026 Kaggle competition discussion forum (analyzed in §6; not redistributed).